

STAT652 Midterm Report

Chun Yin Kong

3/10/2021

Initial Setup

Loading Data and Library

```
library(tictoc)
library(pacman)
p_load(titanic, tidyverse, janitor, naniar, DataExplorer, tidymodels)
data(titanic_train)
data(titanic_test)
set.seed(4)
```

Study the data pattern of both training and testing data set:

```
head(titanic_train)
```

```
##   PassengerId Survived Pclass
## 1           1         0       3
## 2           2         1       1
## 3           3         1       3
## 4           4         1       1
## 5           5         0       3
## 6           6         0       3
##
##                                Name    Sex Age SibSp Parch
## 1                                Braund, Mr. Owen Harris   male  22     1     0
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female  38     1     0
## 3                                Heikkinen, Miss. Laina female  26     0     0
## 4 Futrelle, Mrs. Jacques Heath (Lily May Peel) female    35     1     0
## 5                                Allen, Mr. William Henry   male  35     0     0
## 6                                Moran, Mr. James         male  NA     0     0
##
##      Ticket    Fare Cabin Embarked
## 1    A/5 21171  7.2500      S
## 2    PC 17599 71.2833    C85      C
## 3 STON/O2. 3101282  7.9250      S
## 4    113803 53.1000   C123      S
## 5    373450  8.0500      S
## 6    330877  8.4583      Q
```

```
head(titanic_test)
```

```
##   PassengerId Pclass                                Name    Sex Age
## 1          892      3                                Kelly, Mr. James   male 34.5
## 2          893      3        Wilkes, Mrs. James (Ellen Needs) female 47.0
## 3          894      2                                Myles, Mr. Thomas Francis   male 62.0
## 4          895      3                                Wirz, Mr. Albert   male 27.0
```

```
## 5      896      3 Hirvonen, Mrs. Alexander (Helga E Lindqvist) female 22.0
## 6      897      3      Svensson, Mr. Johan Cervin   male 14.0
## SibSp Parch Ticket   Fare Cabin Embarked
## 1      0      0 330911  7.8292      Q
## 2      1      0 363272  7.0000      S
## 3      0      0 240276  9.6875      Q
## 4      0      0 315154  8.6625      S
## 5      1      1 3101298 12.2875      S
## 6      0      0   7538  9.2250      S
```

Preparing Data for model training

Training and Testing Data

```
titanic_train <- titanic_train %>% clean_names()
#head(titanic_train)
```

Check for duplicate records/examples/rows in your dataset.

```
#get_dupes(titanic_train)
```

```
titanic_train2 <- titanic_train %>% select(-passenger_id, -name, -ticket, -cabin) %>%
  mutate(
    survived = as_factor(survived),
    pclass = as_factor(pclass),
    sex = as_factor(sex),
    embarked = as_factor(embarked)
  )
#head(titanic_train2)
```

```
titanic_test <- titanic_test %>% clean_names()
#head(titanic_test)
```

Check for duplicate records/examples/rows in your dataset.

```
#get_dupes(titanic_test)
```

```
titanic_test2 <- titanic_test %>% select(-passenger_id, -name, -ticket, -cabin) %>%
  mutate(
    pclass = as_factor(pclass),
    sex = as_factor(sex),
    embarked = as_factor(embarked)
  )
#head(titanic_test2)
```

Make the first split with 80% of the data being in the training data set.

```
titanic_train2_split <- initial_split(titanic_train2, prop = 0.8)
#titanic_train2_split
```

Create the recipe for applying the preprocessing. Note the use of `step_nzv()`, which removes any columns that have very low variability, and the use of the `step_meanimpute()` function, which fills in the cells that are missing with the mean of the column.

```
titanic_train2_recipe <- training(titanic_train2_split) %>%
  recipe(survived ~ .) %>%
  step_rm(pclass, sex, embarked) %>%
  step_nzv(all_predictors()) %>%
  step_meanimpute(age) %>%
  prep()

#summary(titanic_train2_recipe)

#tidy(titanic_train2_recipe)
```

Apply the recipe, so the *age* variable should be complete after the imputation.

```
titanic_train2_testing <- titanic_train2_recipe %>%
  bake(testing(titanic_train2_split))

#titanic_train2_testing

titanic_train2_training <- juice(titanic_train2_recipe)

#titanic_train2_training
```

Model Comparison

In this report, I will directly apply code with tuning in order to find the best parameters for each algorithms yielding the best result for model comparison. Then I will pick the model which perform the best to rerun the full titanic_train dataset and produce prediction for the titanic_test dataset.

0. Null Model

Training Model

```
model_null <- logistic_reg(mode = "classification") %>%
  set_engine("glm") %>%
  fit(survived ~ 1, data = titanic_train2_training)
```

Testing Model with Prediction

```
library(yardstick)
pred <- titanic_train2_testing %>%
  select(survived) %>%
  bind_cols(
    predict(model_null, new_data = titanic_train2_testing, type = "class")
  ) %>%
  rename(survived_null = .pred_class)
```

Evaluating the model performance

```
model_null %>%
  predict(titanic_train2_testing) %>%
  bind_cols(titanic_train2_testing) %>%
  conf_mat(truth = survived, estimate = .pred_class)
```

```
##           Truth
## Prediction    0    1
##           0 113  65
##           1   0   0
```

Accuracy and Metrics:

```
result_null_2 <- metrics(pred, survived, survived_null) %>%
  mutate(algorithm = "Null Model")

metrics(pred, survived, survived_null)
```

```
## # A tibble: 2 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 accuracy binary      0.635
## 2 kap     binary        0
```

1. kNN with normalization

Tuning K parameter

In here I use for loop to test the k-nearest neighbor value from 1 to 1000

```
tic()
result_kNN <- data.frame()
for (i in 1:1000){

  kNN_value <- c(i, i)

  titanic_train2a_knn <- nearest_neighbor(neighbors = i) %>%
    set_engine("kkn", scale=TRUE) %>%
    set_mode("classification") %>%
    parsnip::fit(survived ~ ., data = titanic_train2_training)

  result_temp <- titanic_train2a_knn %>%
    predict(titanic_train2_testing) %>%
    bind_cols(titanic_train2_testing) %>%
    metrics(truth = survived, estimate = .pred_class)

  result_temp <- bind_cols(kNN_value, result_temp)


  result_kNN <- bind_rows(result_kNN, result_temp)

}

result_kNN <- result_kNN %>%
  rename(knn_value = ...1)
toc()

## 200.184 sec elapsed

library(ggplot2)
ggplot(data = result_kNN, aes(y = .estimate, x = knn_value, col=.metric)) +
  geom_point() +
  geom_line()
```



Finding the k value yield the best accuracy and kappa:

```
result_kNN %>%
  filter(.estimate == max(.estimate))
```

##	knn_value	.metric	.estimator	.estimate
## 1	172	accuracy	binary	0.7696629
## 2	173	accuracy	binary	0.7696629
## 3	174	accuracy	binary	0.7696629
## 4	175	accuracy	binary	0.7696629
## 5	176	accuracy	binary	0.7696629
## 6	177	accuracy	binary	0.7696629
## 7	178	accuracy	binary	0.7696629
## 8	179	accuracy	binary	0.7696629
## 9	180	accuracy	binary	0.7696629
## 10	181	accuracy	binary	0.7696629
## 11	182	accuracy	binary	0.7696629
## 12	192	accuracy	binary	0.7696629
## 13	193	accuracy	binary	0.7696629
## 14	194	accuracy	binary	0.7696629
## 15	195	accuracy	binary	0.7696629
## 16	241	accuracy	binary	0.7696629
## 17	242	accuracy	binary	0.7696629
## 18	243	accuracy	binary	0.7696629

kNN Model with best k-value

```
result_knn_2 <- result_knn %>%  
  filter(knn_value == "173") %>%  
  mutate(algorithm = "knn")
```

```
result_knn_2
```

```
##   knn_value .metric .estimator .estimate algorithm  
## 1      173 accuracy    binary 0.7696629      knn  
## 2      173      kap     binary 0.4665985      knn
```

2. Boosted C5.0

Tuning Number of Trees in Boosted C5.0

```
result_boosted_c50 <- data.frame()

for (i in 1:100){
  tree_value <- c(i, i)
  titanic_train2a_C50 <- boost_tree(trees = i) %>%
    set_engine("C5.0") %>%
    set_mode("classification") %>%
    fit(survived ~ ., data = titanic_train2_training)


  result_temp <- titanic_train2a_C50 %>%
    predict(titanic_train2_testing) %>%
    bind_cols(titanic_train2_testing) %>%
    metrics(truth = survived, estimate = .pred_class)

  result_temp <- bind_cols(tree_value, result_temp)

  result_boosted_c50 <- bind_rows(result_boosted_c50, result_temp)
}

result_boosted_c50 <- result_boosted_c50 %>%
  rename(tree_value = ...1)

ggplot(data = result_boosted_c50, aes(y = .estimate, x = tree_value, col=.metric)) +
  geom_point() +
  geom_line()
```

Boosted Tree model with best number of tree

Finding the number of trees value yield the best accuracy and kappa:

```
result_boosted_c50 %>%
  filter(.estimate == max(.estimate))

##   tree_value .metric .estimator .estimate
## 1          1 accuracy   binary 0.7752809
## 2          2 accuracy   binary 0.7752809
## 3          3 accuracy   binary 0.7752809
## 4          5 accuracy   binary 0.7752809
## 5          6 accuracy   binary 0.7752809
## 6          7 accuracy   binary 0.7752809

result_boosted_c50_2 <- result_boosted_c50 %>%
  filter(tree_value == "6") %>%
  mutate(algorithm = "Boosted C5.0")

result_boosted_c50_2

##   tree_value .metric .estimator .estimate  algorithm
## 1          6 accuracy   binary 0.7752809 Boosted C5.0
## 2          6      kap    binary 0.4626415 Boosted C5.0
```

3. Random Forest

Tuning Number of Trees in Random Forest

```
result_random_forest <- data.frame()

for (i in 1:1000){
  tree_value <- c(i, i)
  titanic_train2a_ranger <- rand_forest(trees = i) %>%
    set_engine("ranger") %>%
    set_mode("classification") %>%
    fit(survived ~ ., data = titanic_train2_training)


  result_temp <- titanic_train2a_ranger %>%
    predict(titanic_train2_testing) %>%
    bind_cols(titanic_train2_testing) %>%
    metrics(truth = survived, estimate = .pred_class)

  result_temp <- bind_cols(tree_value, result_temp)

  result_random_forest <- bind_rows(result_random_forest, result_temp)
}

result_random_forest <- result_random_forest %>%
  rename(tree_value = ...1)

ggplot(data = result_random_forest, aes(y = .estimate, x = tree_value, col=.metric)) +
  geom_point() +
  geom_line()
```



Finding the number of trees value yield the best accuracy and kappa:

```
result_random_forest %>%
  filter(.estimate == max(.estimate))

##   tree_value .metric .estimator .estimate
## 1         13 accuracy   binary 0.7808989
## 2         40 accuracy   binary 0.7808989
## 3         61 accuracy   binary 0.7808989
## 4         85 accuracy   binary 0.7808989
## 5        125 accuracy   binary 0.7808989
## 6        202 accuracy   binary 0.7808989

result_random_forest_2 <- result_random_forest %>%
  filter(tree_value == "85") %>%
  mutate(algorithm = "Random Forest")

result_random_forest_2

##   tree_value .metric .estimator .estimate   algorithm
## 1         85 accuracy   binary 0.7808989 Random Forest
## 2         85      kap    binary 0.5258844 Random Forest
```

4. Logistic Regression using regularization

Tuning penalty and mixture

```
library(discrim)
result_log_reg <- data.frame()

for (j in seq(0, 1, by=0.1)){
  for (i in seq(0, 0.01, by=0.001)){
    penalty_val <- c(i, i)
    mixture_val <- c(j, j)
    titanic_train2a_glm <- logistic_reg(penalty = i, mixture = j) %>%
      set_engine("glmnet") %>%
      set_mode("classification") %>%
      fit(survived ~ ., data = titanic_train2_training)

    result_temp <- titanic_train2a_glm %>%
      predict(titanic_train2_testing) %>%
      bind_cols(titanic_train2_testing) %>%
      metrics(truth = survived, estimate = .pred_class)


    result_temp <- bind_cols(penalty_val, mixture_val, result_temp)

    result_log_reg <- bind_rows(result_log_reg, result_temp)
  }
}

result_log_reg <- result_log_reg %>%
  rename(penalty_val = ...1) %>%
  rename(mixture_val = ...2) %>%
  mutate(mixture_val = as.factor(mixture_val))

result_log_reg1a <- result_log_reg %>%
  filter(.metric == "accuracy")

ggplot(data = result_log_reg1a, aes(y = .estimate, x = penalty_val, col=mixture_val)) +
  geom_point() +
  geom_line()
```



From the above short comparison, we find that although changing the penalty and mixture value of the logistic regression will affect the accuracy, but the difference in accuracy value is so small that it is not so significant in changing the penalty and mixture in this data set. Hence below is a random pick of penalty to produce the accuracy result. For the value of mixture (L1 Regularization), it will change when it is greater than 0.5 or less than 0.5 for each penalty value. i.e. the tuning point of accuracy of each penalty value is in mixture value = 0.5. Hence I set the mixture value to 0.5 to give the highest accuracy for the logistic regression prediction.

```
result_log_reg_2 <- logistic_reg(penalty = 0.001, mixture = 0.5) %>%
  set_engine("glmnet") %>%
  set_mode("classification") %>%
  fit(survived ~ ., data = titanic_train2_training) %>%
  predict(titanic_train2_testing) %>%
  bind_cols(titanic_train2_testing) %>%
  metrics(truth = survived, estimate = .pred_class) %>%
  mutate(penalty_val = 0.001,
         mixture_val = 0.5,
         algorithm = "Logistic Regression with Regularization")

result_log_reg_2
```

```
## # A tibble: 2 x 6
##   .metric .estimator .estimate penalty_val mixture_val algorithm
##   <chr>    <chr>         <dbl>    <dbl>    <dbl> <chr>
## 1 accuracy binary         0.747      0.001      0.5 Logistic Regression wit~
## 2 kap     binary         0.375      0.001      0.5 Logistic Regression wit~
```

5. Naive Bayes

Tuning Laplace Estimator

```
library(discrim)
result_naive_bayes <- data.frame()

for (i in 0:999){
  laplace_est <- c(i, i)
  titanic_train2a_nb <- naive_Bayes(Laplace = i) %>%
    set_engine("klaR") %>%
    set_mode("classification") %>%
    fit(survived ~ ., data = titanic_train2_training)


  result_temp <- titanic_train2a_nb %>%
    predict(titanic_train2_testing) %>%
    bind_cols(titanic_train2_testing) %>%
    metrics(truth = survived, estimate = .pred_class)

  result_temp <- bind_cols(laplace_est, result_temp)

  result_naive_bayes <- bind_rows(result_naive_bayes, result_temp)
}

result_naive_bayes <- result_naive_bayes %>%
  rename(laplace_est = ...1)

ggplot(data = result_naive_bayes, aes(y = .estimate, x = laplace_est, col=.metric)) +
  geom_point() +
  geom_line()
```



From the graph above we can see that the accuracy doesn't change when the Laplace Estimator varies. Hence in this model I will pick the Laplace Estimator of 0 for the best parameter for this model, that giving best accuracy.

```
result_naive_bayes_2 <- result_naive_bayes %>%
  filter(laplace_est == "0") %>%
  mutate(algorithm = "Naive Bayes")
```

```
result_naive_bayes_2
```

```
##   laplace_est .metric .estimator .estimate algorithm
## 1           0 accuracy   binary 0.7247191 Naive Bayes
## 2           0      kap    binary 0.3756621 Naive Bayes
```

Comparing all models

```
result_all_models <- bind_rows(result_null_2, result_knn_2, result_boosted_c50_2, result_random_forest_2)
result_all_models
```

```
## # A tibble: 12 x 9
##   .metric .estimator .estimate algorithm knn_value tree_value penalty_val
##   <chr>    <chr>      <dbl> <chr>      <int>    <int>      <dbl>
## 1 accuracy binary      0.635 Null Model      NA        NA        NA
## 2 kap      binary      0      Null Model      NA        NA        NA
## 3 accuracy binary      0.770 kNN          173       NA        NA
## 4 kap      binary      0.467 kNN          173       NA        NA
## 5 accuracy binary      0.775 Boosted C5.0    NA         6        NA
## 6 kap      binary      0.463 Boosted C5.0    NA         6        NA
```

```
## 7 accuracy binary      0.781 Random Forest      NA      85      NA
## 8 kap      binary      0.526 Random Forest      NA      85      NA
## 9 accuracy binary      0.747 Logistic Regr~      NA      NA      0.001
## 10 kap      binary      0.375 Logistic Regr~      NA      NA      0.001
## 11 accuracy binary      0.725 Naive Bayes      NA      NA      NA
## 12 kap      binary      0.376 Naive Bayes      NA      NA      NA
## # ... with 2 more variables: mixture_val <dbl>, laplace_est <int>
```

From the table above we see that after running all models, random forest algorithm gives the highest accuracy in the titanic data set. Hence in the following pages I will use random forest to rerun the full titanic data set and for model evaluation.

ROC Curves

```
null_roc <- model_null %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)

knn_roc <- titanic_train2a_knn %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)

c50_roc <- titanic_train2a_C50 %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)


rf_roc <- titanic_train2a_ranger %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)

glm_roc <- titanic_train2a_glm %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)

nb_roc <- titanic_train2a_nb %>%
  predict(titanic_train2_testing, type = "prob") %>%
  bind_cols(titanic_train2_testing) %>%
  roc_curve(survived, .pred_0)
```

ROC Curve for Best Model - Random Forest

```
autoplot(rf_roc)
```

Rerun Full Titanic Dataset with Best Model

```
titanic_train3 <- titanic_train2 %>%
  na.omit()

titanic_test3 <- titanic_test2 %>%
  na.omit()

titanic_test4 <- titanic_test %>%
  na.omit() %>%
  select(passenger_id)

titanic_rf_full_model <- rand_forest(trees = 85) %>%
  set_engine("ranger") %>%
  set_mode("classification") %>%
  fit(survived ~ ., data = titanic_train3)

result_rf_full <- titanic_rf_full_model %>%
  predict(titanic_test3) %>%
  bind_cols(titanic_test4)

result_rf_full <- result_rf_full[, c(2, 1)]
result_rf_full <- result_rf_full %>%
  rename(PassengerId = passenger_id,
         Survived = .pred_class)

head(result_rf_full)

## # A tibble: 6 x 2
##   PassengerId Survived
##         <int> <fct>
## 1         892 0
## 2         893 0
## 3         894 0
## 4         895 0
## 5         896 0
## 6         897 0

write.csv(result_rf_full, "titanic_full_prediction.csv")
```